



Traffic Sign Recognition

BY:

Ayham Zaza

For

SDCND Program

April. 2017



Traffic Sign Recognition

OBJECTIVES

- Load the data set.
- Explore, summarize and visualize the data set
- Design, train and test a model architecture
- Use the model to make predictions on new images
- Analyze the softmax probabilities of the new images
- Summarize the results with a written report

RUBRIC POINTS

Here I will consider the [rubric points](#) individually and describe how I addressed each point in my implementation.

WRITEUP / README

1. **Provide a Writeup / README that includes all the rubric points and how you addressed each one. You can submit your writeup as markdown or pdf. You can use this template as a guide for writing the report. The submission includes the project code.**

You're reading it! and here is a link to my [project code](#)

Data Set Summary & Exploration

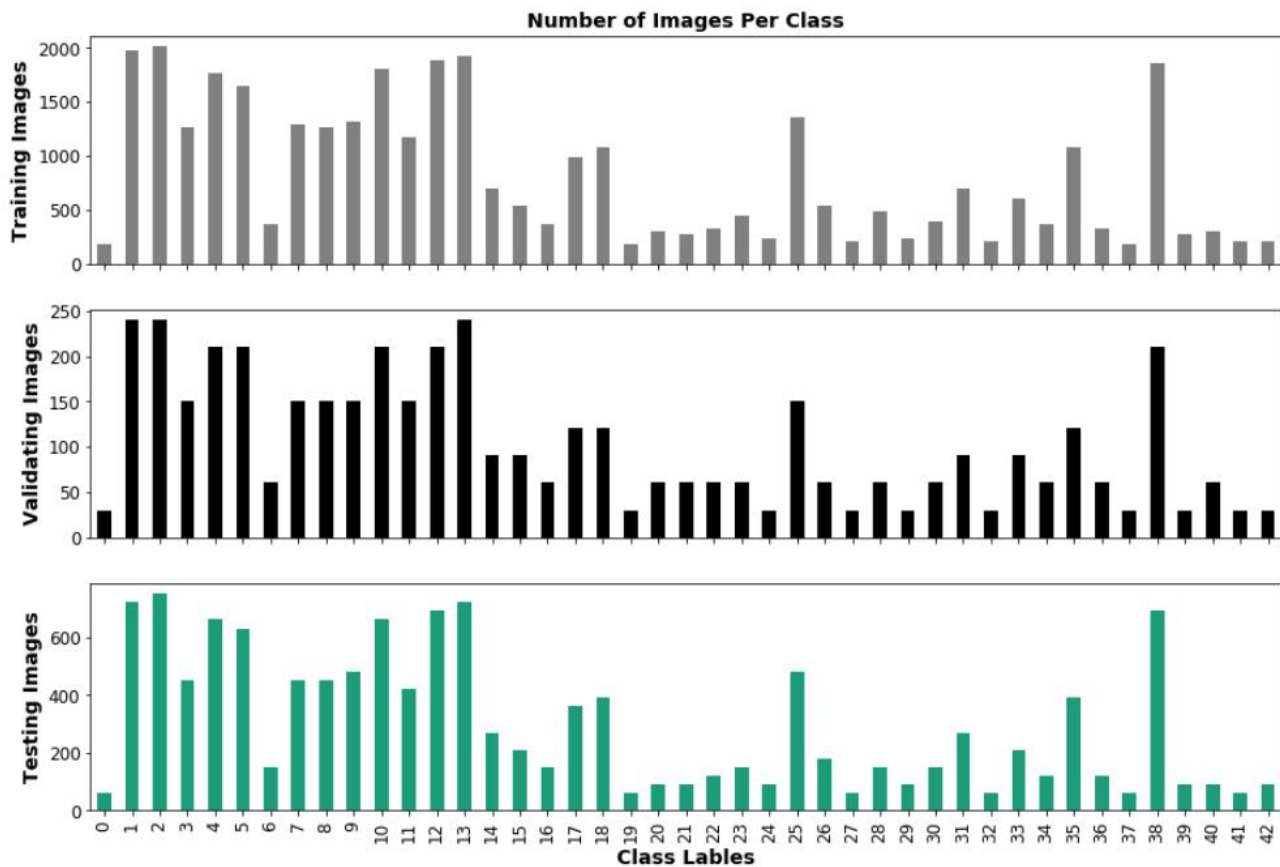
1. **Provide a basic summary of the data set. In the code, the analysis should be done using python, numpy and/or pandas methods rather than hardcoding results manually.**

I used the pandas library to calculate summary statistics of the traffic signs data set:

- The size of training set is **34799**
- The size of the validation set is **4410**
- The size of test set is **12630**
- The shape of a traffic sign image is **32x32x3**
- The number of unique classes/labels in the data set is **43**

2. **Include an exploratory visualization of the dataset.**

Here is an exploratory visualization of the data set. It is a bar chart showing how the data ...



Design and Test a Model Architecture

1. Describe how you preprocessed the image data. What techniques were chosen and why did you choose these techniques? Consider including images showing the output of each preprocessing technique. Pre-processing refers to techniques such as converting to grayscale, normalization, etc. (OPTIONAL: As descry bed in the "Stand Out Suggestions" part of the rubric, if you generated additional data for training, describe why you decided to generate additional data, how you generated the data, and provide example images of the additional data. Then describe the characteristics of the augmented training set like number of images in the set, number of images for each class, etc.)

As could be seen from the above figure, the distribution of class labels is imbalanced. This will necessarily cause some biasing when training the network by giving more weights to certain images. Thus, additional images were adaptively added to every class as needed to compensate for such discrepancy. The number of additional images was heuristically determined. Two functions served that purpose

- `adjust_images()`:

defines 13 transformations including image rotation, flipping, smoothing ..etc and returns a number of newly generated images as specified by the pivot.

- **generate_images()**:

generate enough images for every class as detected and required using **get_statistics()**.

I decided to convert the images to grayscale because to reduce image dimensions as well as any additional unnecessary information. This will speed up the training process.

I also normalized the image data because reduce the gap among various image data and to drive faster training.

Here is an example of a traffic randomly selected traffic sign images



The size of all the sets were verified after that addition operation and summarized below along with the new bar plot that shows the new distribution.

For Training Set

The length before extendeing images = 34799 and the length of the assoc. label = 34799

The length after extendeing images = 39418 and the length of the assoc. label = 39418

For Validation Set


The length before extendeing images = 4410 and the length of the assoc. label = 4410

The length after extendeing images = 4980 and the length of the assoc. label = 4980

For Testing Set

The length before extendeing images = 12630 and the length of the assoc. label = 12630

The length after extendeing images = 15450 and the length of the assoc. label = 15450



and here is a snapshot of some randomly selected validation images after augmentation.

0: Stop



1: Dangerous curve to the left



2: Road work



3: Speed limit (50km/h)



4: Go straight or right




5: No passing for vehicles over 3.5 metric tons



6: Double curve



7: Ahead only



8: Ahead only




9: Turn left ahead



10: Speed limit (120km/h)



11: Beware of ice/snow



12: Speed limit (70km/h)




13: No passing for vehicles over 3.5 metric tons



14: Double curve



15: Speed limit (60km/h)



3. Describe what your final model architecture looks like including model type, layers, layer sizes, connectivity, etc.) Consider including a diagram and/or table describing the final model.

Training Environment

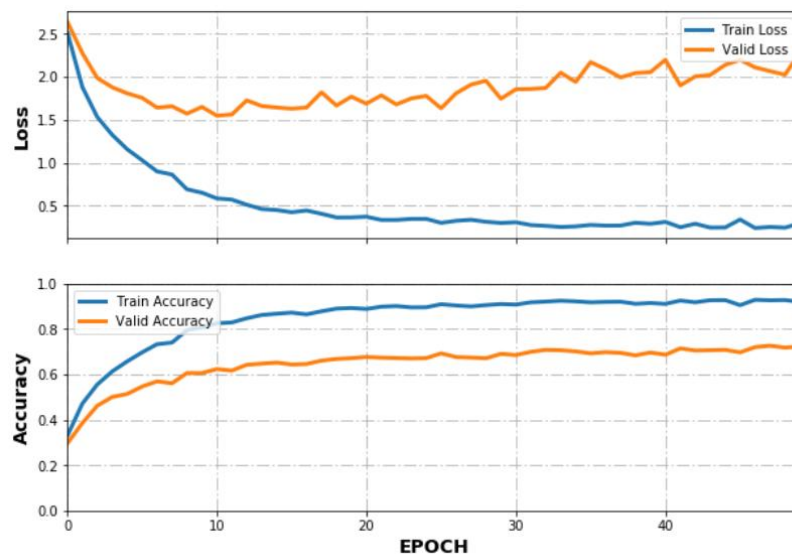
Laptop with windows 10 operating system and x64- based processor, i7 architecture with 2.2 GHz clock. 16 GB memory and NVIDIA GeForce GT 750M GPU.

Various models have been tried and tested, here I list those trials as well as some comments. For all models the input is **32x32x1**

Model 1: Adapted From LeNET

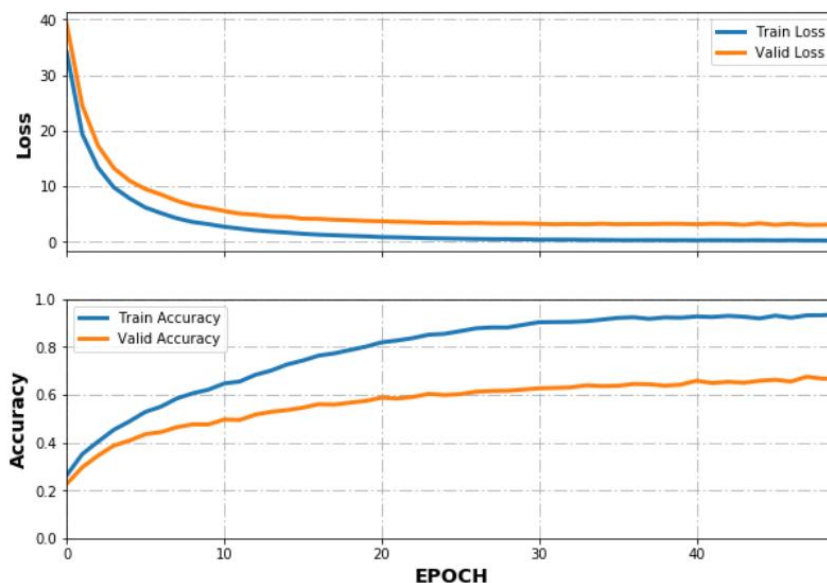
| Layer | Output | characteristics |
|---------------------|-----------|---------------------------|
| 32x32x1 Input Image | | |
| 5x5 Convolution | 28x28x8 | 1x1 stride, valid padding |
| RELU | | |
| Max Pooling | 14x14x8 | 2x2 stride |
| 5x5 Convolution | 10x10x16 | 1x1 stride, valid padding |
| RELU | | |
| Max Pooling | 5x5x16 | 2x2 stride |
| Fully Connected | 400 x2000 | |
| Fully Connected | 2000x256 | |
| Output | 256x43 | |

The following figure shows the accuracy and loss output for both training and validation data after **50 epochs**, using **0.0001** training rate and **AdamOptimizer**



Not only this model training accuracy seems to plateau at 92% but also **overfits** the data. There is no point in examining dropout or any other regularization techniques over this model.


Next and given the same network structure (2 conv-nets followed by 2 fully connected), I increased the depth in both convolution layers to **32** and **64** respectively and adjusted all other parameters. The following plot was obtained for the same testing conditions as above:



This adjustment produced a smoother loss curve for the validation data. The learning rate seems very low. The following suggestions will be considered for other models:

1. Experimenting with lower learning rates
2. Having large depth for conv-nets
3. Experimenting with other solvers

Following the previous suggestions and focusing on item 1 & 3, I tried **AdagradOptimizer** with **0.01** Learning rate and keeping everything the same. The following plot was obtained:



The training time took around 25 minutes with the following final output:


```
EPOCH 50 ...  
Train loss = 0.164  
Train accuracy = 0.960  
Valid loss = 1.035  
Valid accuracy = 0.733  
Time: 1487.979 seconds
```

The Following could be concluded:

1. The smoothness in both learning and validation curves suggests that higher epochs may indeed reach higher levels of accuracy.
2. Dropout (and/or) regularization could be utilized to enhance the validation accuracy.

By adjusting keep probability in **dropout** to 0.77, and lower the number of hidden units by half the following was obtained:

```
EPOCH 70 ...  
Train loss = 0.217  
Train accuracy = 0.944  
Valid loss = 0.834  
Valid accuracy = 0.755  
Time: 1460.281 seconds
```



The above results and the slow convergence rate suggest the following:


- Increase the number of epoch to a higher number. 500 for instance, while increasing the learning rate. I cannot afford doing this solution on my laptop. Already with 50 epochs lot of heat is generated.
- Changing the model by reducing the number of nodes in the hidden layers and adding more layer
- Checking the effect of adding an inception layer.

Model 2: Increasing the depth of the Network

| Layer | Output | characteristics |
|------------------------|-----------------|---------------------------|
| 32x32x1 Input Image | | |
| 3x3 Convolution + ReLU | 30x30x16 | 1x1 stride, valid padding |
| Max Pooling | 15x15x16 | 2x2 stride |
| 4x4 Convolution + ReLU | 12x12x24 | 1x1 stride, valid padding |
| Max Pooling | 6x6x24 | 2x2 stride |
| 1x1 Convolution + ReLU | 6x6x32 | 1x1 stride, SAME padding |
| Max Pooling | 3x3x3x32=864 ?? | 2x2 stride |
| Fully Connected | 864 x 600 | |
| Fully Connected | 600x 256 | |
| Fully Connected | 256 x 128 | |
| Fully Connected | 128 x 96 | |
| Output | 96 x 43 | |


The following output was obtained at epoch 70

```
EPOCH 70 ...  
Train loss = 0.082  
Train accuracy = 0.981  
Valid loss = 1.356  
Valid accuracy = 0.709  
Time: 1257.198 seconds
```




The learning rate of **0.01** seems just fine and **AdagradOptimizer** produced a very high training accuracy that is smooth and monotonically decreasing. To cope with the degradation in validation accuracy, a drop out layer was introduced after the first and second convolution layers. Keep probability was chosen to be 0.5. The following was obtained:

```
EPOCH 100 ...  
Train loss = 1.787  
Train accuracy = 0.475  
Valid loss = 2.417  
Valid accuracy = 0.326  
Time: 4036.159 seconds
```




The previous figure suggests that the use of drop out is already exaggerated. In the next experiment a drop out of 0.7 was tried after applying it to the first convolution layer.

```
EPOCH 100 ...  
Train loss = 0.102  
Train accuracy = 0.974  
Valid loss = 1.239  
Valid accuracy = 0.696  
Time: 2376.808 seconds
```



Next I tried lower the same drop out to as low as 0.3 and the following was obtained:

```
EPOCH 100 ...  
Train loss = 1.108  
Train accuracy = 0.710  
Valid loss = 1.779  
Valid accuracy = 0.467  
Time: 8535.246 seconds
```



This may lead to good results at the long run, but I cannot try running more than 100 epochs on my laptop as the excessive generated heat may damage the computer. Therefore, another architecture is to be examined with the following suggestions:

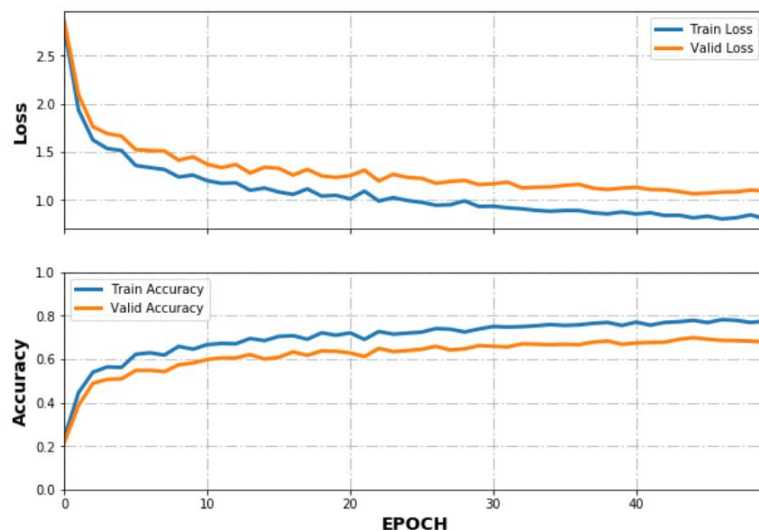
- Lowering the number of nodes at each layer.
- Reducing the depth of the network
- Trying the inception model and average lower pooling

Model 3: Examining with the inception model

| Layer | Output | characteristics |
|------------------------|----------|---------------------------|
| 32x32x1 Input Image | | |
| 3x3 Convolution + RELU | 30x30x4 | 1x1 stride, valid padding |
| 3x3 Convolution + RELU | 28x28x8 | 1x1 stride, valid padding |
| 5x5 Convolution + RELU | 24x24x12 | 1x1 stride, valid padding |
| Max Pooling | 12x12x12 | 2x2 stride |
| Inception Layer | | |
| 1x1 Convolution + ReLU | 12x12x32 | 1x1 stride, SAME padding |
| 3x3 Convolution + ReLU | 12x12x32 | 1x1 stride, SAME padding |
| 5x5 Convolution + ReLU | 12x12x32 | 1x1 stride, SAME padding |
| Average Pooling | 4x4x32 | 3x3 stride |
| Fully Connected | 512 x256 | |
| Fully Connected | 256 x 94 | |
| Output | 94 x 43 | |

AdagradOptimizer with training rate of 0.01 was utilized here on the extended list of gray images without normalization. The following output was obtained at epoch 50

```
EPOCH 50 ...  
Train loss = 0.808  
Train accuracy = 0.775  
Valid loss = 1.096  
Valid accuracy = 0.680  
Time: 1733.895 seconds
```



The behavior of the validation rate follows nicely the one for the learning rate. However, they don't provide good accuracy levels at epoch 50. L1 regularization was adapted here.

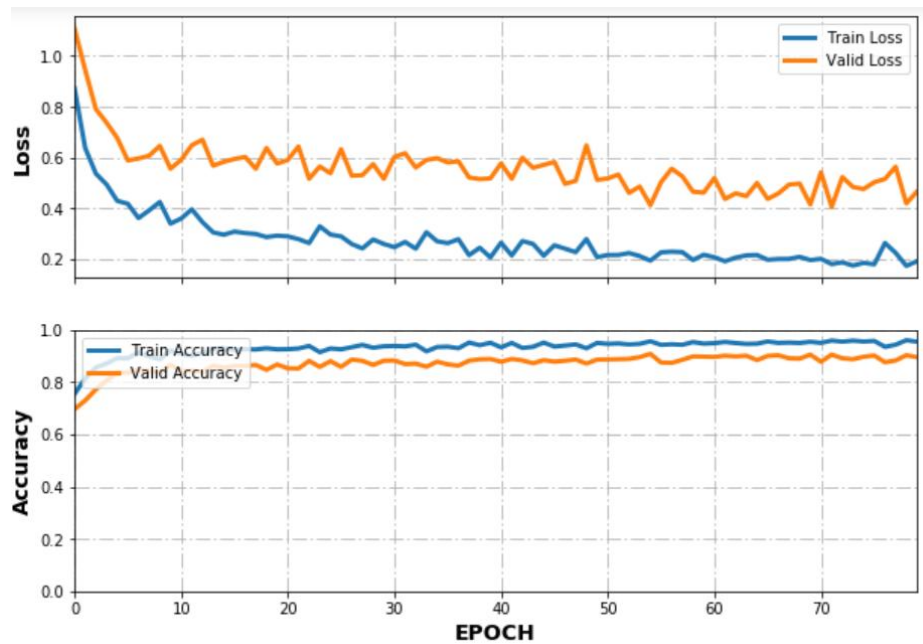
Model 4: Combination of the inception model

Lot and lot of experiments have been tried out using a combination of different models. What was interesting however, was the following experiment where I have not added any extra images and utilized the color images directly. Because of certain stability issues, I decided to go for L1 regularization with AdamOptimizer and 0.001 learning rate.

| Layer | Output | characteristics |
|------------------------|-----------|---------------------------|
| 32x32x1 Input Image | | |
| 3x3 Convolution | 30x30x4 | 1x1 stride, valid padding |
| Inception Layer | | |
| 1x1 Convolution + ReLU | 30x30x8 | 1x1 stride, SAME padding |
| 3x3 Convolution + ReLU | 30x30x8 | 1x1 stride, SAME padding |
| 5x5 Convolution + ReLU | 30x30x8 | 1x1 stride, SAME padding |
| Average Pooling | 15x15x8 | 2x2 stride |
| Inception Layer | | |
| 1x1 Convolution + ReLU | 15x15x8 | 1x1 stride, SAME padding |
| 3x3 Convolution + ReLU | 15x15x8 | 1x1 stride, SAME padding |
| 5x5 Convolution + ReLU | 15x15x8 | 1x1 stride, SAME padding |
| Average Pooling | 7x7x8 | 2x2 stride |
| 3x3 Convolution | 5x5x16 | 1x1 stride, valid padding |
| Fully Connected | 400 x 400 | |
| Fully Connected | 400 x 86 | |
| Output | 86 x 43 | |

Next a combination of regularization and drop out values are examined. The following is the best that could be achieved to the best of my educated guess.

```
EPOCH 80 ...  
Train loss = 0.192  
Train accuracy = 0.955  
Valid loss = 0.466  
Valid accuracy = 0.896  
Time: 1824.964 seconds
```



Clearly, this model, performs much better than the previous ones. It could be further smoothed by creating balanced colored classes in the training set. Also drop out may be applied to push the learning. In addition, if the learning rate is lowered, better results may be achieved.

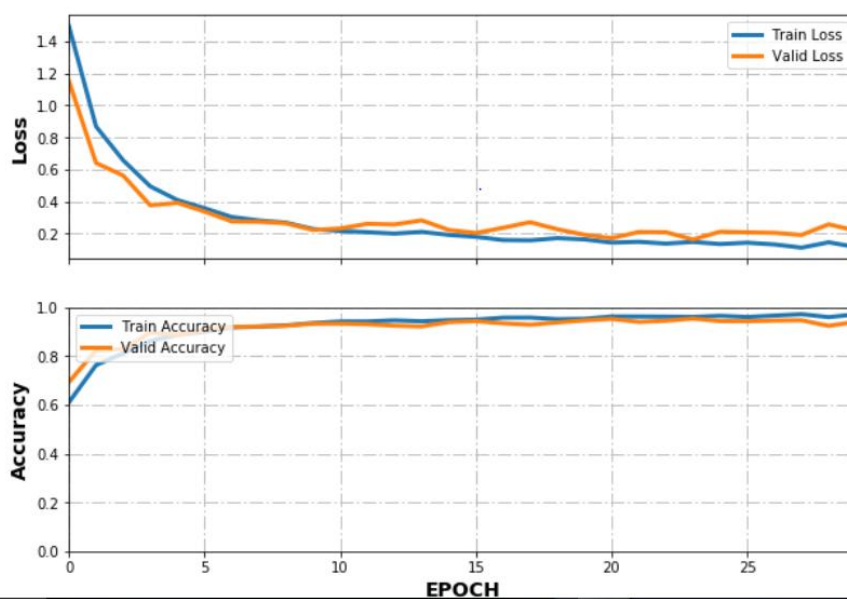
FINAL MODEL

Model 5: Simple Conv Layers and L2 Regularization

Focusing back on gray scale images but this time by applying L2 regularization. The results obtained achieves the required validation accuracy. The learning rate was 0.0004 and the optimization was done using AdamOptimizer.

| Layer | Output | characteristics |
|------------------------|-----------|---------------------------|
| 32x32x1 Input Image | | |
| 3x3 Convolution + ReLU | 30x30x8 | 1x1 stride, valid padding |
| Average Pooling | | k =2 |
| 4x4 Convolution + ReLU | 15x15x16 | 1x1 stride, valid padding |
| Average Pooling | | k =1 |
| 2x2 Convolution + ReLU | 12x12x24 | 1x1 stride, valid padding |
| Average Pooling | | k =2 |
| Fully Connected + ReLU | 600x600 | |
| Fully Connected + ReLU | 600 x 215 | |
| Output Connected | 215 x 43 | |

```
EPOCH 30 ...  
Train loss = 0.117  
Train accuracy = 0.972  
Valid loss = 0.219  
Valid accuracy = 0.941  
Time: 442.872 seconds
```



The above model can surely be enhanced further by trying different optimizers and various training rates.

4. Describe how you trained your model. The discussion can include the type of optimizer, the batch size, number of epochs and any hyperparameters such as learning rate.

To details were highlighted in the previous questions.

- The learning rate was 0.0004,
- Epochs: 30
- Batch Size: 128
- Optimizer: AdamOptimizer
- Loss Operation: Cross Entropy with L2 regularization

5. Describe the approach taken for finding a solution and getting the validation set accuracy to be at least 0.93. Include in the discussion the results on the training, validation and test sets and where in the code these were calculated. Your approach may have been an iterative process, in which case, outline the steps you took to get to the final solution and why you chose those steps. Perhaps your solution involved an already well known implementation or architecture. In this case, discuss why you think the architecture is suitable for the current problem.

The answer to this question was also detailed when answering Q3. Here I reemphasize some points:

- training set accuracy: **0.972**
- validation set accuracy: **0.941**
- test set accuracy: **0.911**

If an iterative approach was chosen:

- I started with LeNET architecture and built things from there. For this problem, that architecture underfits the data, so I had to add extra layer and adjusted several parameters like the pooling size and others. The details are previously detailed
- The convolution layers summarize values of a group of pixels. They therefore, reduce the complexity of the system and helps in better classifications. Drop out is somehow useful as it forces the network to learn about certain parameters in the absence of others. Nevertheless, placing them at a perfect position is challenging.

Test a Model on New Images

1. Choose five German traffic signs found on the web and provide them in the report. For each image, discuss what quality or qualities might be difficult to classify.

Here are six German traffic signs that I found on the web:



The first image might be difficult to classify because the number may be mistaken with other possible speeds like 20 or 90. The last image could also be difficult to classify because it might look as the caution sign, image 2. Image 5 is challenging as well because the sign is tilted.

2. Discuss the model's predictions on these new traffic signs and compare the results to predicting on the test set. At a minimum, discuss what the predictions were, the accuracy on these new predictions, and compare the accuracy to the accuracy on the test set (OPTIONAL: Discuss the results in more detail as described in the "Stand Out Suggestions" part of the rubric).

Here are the results of the prediction:




The model could correctly guess 5 of the 6 traffic signs, which gives an accuracy of 83.3%. This is in good accordance with the test accuracy.


Describe how certain the model is when predicting on each of the five new images by looking at the softmax probabilities for each prediction. Provide the top 5 softmax probabilities for each image along with the sign type of each probability. (OPTIONAL: as described in the "Stand Out Suggestions" part of the rubric, visualizations can also be provided such as bar charts)

The model is certain when predicting the General Caution Sign with probability 1. It is almost sure when predicting dangerous curve, straight or left and the traffic sign with probability > 0.9 . it favours 0.6 of the time and mispredict speed limit 30 with 20. The following table summarizes the results.


Speed limit (20km/h): 0.949
 Keep left: 0.042
 Speed limit (30km/h): 0.005
 End of speed limit (80km/h): 0.001
 Speed limit (50km/h): 0.001




General caution: 1.000
 Pedestrians: 0.000
 Traffic signals: 0.000
 Right-of-way at the next intersection: 0.000
 Road narrows on the right: 0.000




Dangerous curve to the right: 0.935
 Dangerous curve to the left: 0.065
 Children crossing: 0.000
 Slippery road: 0.000
 Beware of ice/snow: 0.000




Go straight or left: 0.853
 Go straight or right: 0.146
 Turn right ahead: 0.000
 Keep left: 0.000
 Speed limit (20km/h): 0.000



Stop: 0.613
 Speed limit (20km/h): 0.116
 Go straight or left: 0.084
 Turn left ahead: 0.057
 Turn right ahead: 0.038



Traffic signals: 0.947
 General caution: 0.051
 Road narrows on the right: 0.001
 Pedestrians: 0.000
 Dangerous curve to the left: 0.000



Conclusion

This project was a chance to strengthen my knowledge about deep neural network in general and convolutional variation in particular. I have already appreciated the importance of experience when trying out so many variations to come up with acceptable accuracy levels in the end. Small variations in some occasions caused the network to either fluctuate, or saturate. I also experimented various famous models.

In the end, I would like to thanks for all the exerted efforts by the Udacity team who do not save any effort to make this track enjoyable.

(Optional) Visualizing the Neural Network (See Step 4 of the Ipython notebook for more details)

Discuss the visual output of your trained network's feature maps. What characteristics did the neural network use to make classifications?

To be continued in future work